

— NEXTUP SOCAL · VIRTUAL EVENT

Reducing AI Operating Costs

What the tokens actually cost, why the invoice grows faster than the usage, and the five levers that bring it back down.



upskilled.consulting/nextup-socal/

Grab the slides

Nick McCarty

Upskilled Consulting

Fri Aug 21, 2026 · 12:00 PM PST

Part two of the NextUp SoCal AI series

In June we mapped the levels. Today: what each one costs to run.



- ✓ **Level 1 – Assistant.** One prompt, one answer. Cheap enough that nobody checks.
- ✓ **Level 2 – Collaborator.** The tool knows your context. More reading, more tokens.
- ✓ **Level 3 – Chief of Staff.** It works without you. It reads *constantly*.

WHERE TODAY PICKS UP

74% of companies plan to deploy agentic AI within two years. So the useful question now is a budgeting one: what does Level 3 cost to run for a year?

From “it’s cheap” to a number you can budget.

01 · THE UNIT

What is a token, and what does one actually cost?

Pages, meetings, emails — converted into the thing you’re billed for.

By the end of it you can price any document your team touches.

02 · THE MULTIPLIER

Why a fraction of a cent becomes half a million dollars.

Three scenarios, same 600 people, same price list, 47× apart.

Same tool, same headcount — three very different invoices.

03 · THE LEVERS

Five ways to cut the bill without cutting the capability.

- Right-size the model
- Script the plain steps
- Engineer the context
- Stop redoing work
- Run commodity work local

What you **don't** spend is also a lever

What you keep out of the prompt, and what you keep out of your stack.

And the cost of standing still

The most expensive line item is usually the work you didn’t automate — and who in the room never got the tool.

You're not billed per prompt. You're billed **per token**.

THE QUESTION

Sum mar ize the Q 3 category review deck for the leadership team .

THE ANSWER

Category volume rose 4.1% against a 2.3% plan, but nearly all of the gain came from promotional units — base velocity was flat. Gross margin fell 60 basis points, driven by deeper promo depth in weeks 7 through 11 and a mix shift toward the value tier. Two SKUs account for 71% of the margin decline; both ran continuous promotion all quarter. Recommend holding promo depth flat through the December reset, moving those two SKUs to an every-other-week cadence, and revisiting value-tier facings at the January line review.

TOKENS IN

14

the question you asked

TOKENS OUT

180

the summary it wrote back

COST OF THE QUESTION

\$0.00002

14 tokens @ \$1.25 / 1M

COST OF THE ANSWER

\$0.0018


180 tokens @ \$10 / 1M

Output is priced **8× higher** than input — \$10 vs \$1.25 per million tokens. On this request **99% of the cost is the answer**. Agents flip that completely: they read far more than they write.

Your work, converted into the billable unit.

Every artifact your team touches has a token weight.

ONE THOUSAND TOKENS



≈ 750 words

≈ 4,000 characters

≈ 1½ pages of text

WHAT YOUR ARTIFACTS WEIGH, IN TOKENS



70× the meeting – and, as you’re about to see, still under twenty cents.


Here is what one of everything costs.

Read it in, hand a summary back. Common artifacts, priced.

TOTAL COST PER ITEM — READ IN, SUMMARY BACK OUT

GPT-5 STANDARD

SMALL-MODEL TIER (~10× CHEAPER)




Nothing here is expensive — which is why the invoice often surprises people as they move up the levels.

For perspective: what everyday work costs.

A single email	< \$0.001
A 10-page report	≈ 2¢
A one-hour meeting	≈ 2¢
A 100-page report	≈ 16¢
A day of email — 50 messages	≈ 5¢
1,000 reports × 100 pages each	≈ \$158

1,000 REPORTS · 100 PAGES EACH



\$157.50

READ AND SUMMARIZED, FOR ABOUT THE PRICE OF ONE TEAM LUNCH

The unit is cheap. The bill is a volume problem — and volume is a decision we have influence over.

— 02 · THE MULTIPLIER

Nobody's budget breaks on price.
It breaks on **volume**.

$$\begin{array}{ccccccc} \$0.02 & \times & 40 & \times & 600 & \times & 250 & = & \$120,000 \\ \text{COST PER TASK} & & \text{TASKS PER DAY} & & \text{PEOPLE} & & \text{WORKING DAYS} & & \text{A YEAR} \end{array}$$

The assistant. Chat, summarize, draft.

PER EMPLOYEE, PER DAY

INPUT — WHAT IT READS

20,000
tokens


OUTPUT — WHAT IT WRITES

5,000
tokens

AT GPT-5 STANDARD PRICING

\$1.25 per 1M input tokens
\$10 per 1M output tokens

WHERE THE DAILY COST GOES



DAILY, PER EMPLOYEE

\$0.075

× 600 employees × 250 working days

EVERY YEAR

The agent.

Agents read far more than they write.

PER EMPLOYEE, PER DAY

INPUT — WHAT IT READS

500,000

tokens

OUTPUT — WHAT IT WRITES

20,000

tokens

AT GPT-5 STANDARD PRICING

\$1.25 per 1M input tokens

\$10 per 1M output tokens

WHERE THE DAILY COST GOES



input \$0.625 · 76%

output \$0.200 · 24%

DAILY, PER EMPLOYEE

\$0.825

× 600 employees × 250 working days

EVERY YEAR

11× scenario 1 — same 600 people

Always-on automation. Agents that never stop watching.

PER EMPLOYEE, PER DAY

INPUT — WHAT IT READS

2,000,000

tokens

OUTPUT — WHAT IT WRITES

100,000

tokens

AT GPT-5 STANDARD PRICING

\$1.25 per 1M input tokens

\$10 per 1M output tokens

WHERE THE DAILY COST GOES



input \$2.50 · 71%

output \$1.00 · 29%

DAILY, PER EMPLOYEE

\$3.50

× 600 employees × 250 working days


EVERY YEAR

47× scenario 1 — same 600 people

**Same 600 people. Same 250 days.
Same price list. 47× apart.**



Five levers. Same capability, a fraction of the bill.



We're paying frontier prices for a 1.7% edge.

BENCHMARK SCORE — BEST PAID vs BEST OPEN-WEIGHT

Best paid model

premium per-token pricing



Best open-weight model


run it at a fraction of the cost



8% → 1.7%

performance gap, closed in a single year

OF 100 TASKS...



- 80 routine — cheap tier
- 20 genuinely hard — frontier

RELATIVE SPEND ON THE SAME 100 TASKS

Everything on the frontier model



Routed — 80 cheap / 20 frontier



$$(20 \times 1.0) + (80 \times 0.1) = 28$$

-72%

on the same work, from a routing rule

Defaulting to the premium model for every task is a tax on habit — not a quality decision.

Most of what you're automating doesn't need to think.

ANNUAL COST



The test is simple: if you can write the rule down, you do not need to rent judgment to follow it.

Where the task genuinely needs judgment or language, the model earns its price. Spend it there.

Your agent is re-reading the same page all day long.

TOTAL READING THE AGENT HAS PAID FOR, IN TOKENS




Cap what the agent re-reads and the same work costs about a quarter as much.*

* Assuming two-thirds of tokens never reach the finished answer.

The chat window is a shredder with a very good memory of nothing.

60 CHAT SESSIONS THIS MONTH



- 18 of them re-answered something already answered
- and none of the answers were saved anywhere

AT 30% REWORK, ON A \$525,000 BILL

\$157,500

a year, re-answering answered questions
— before you count anyone's salary time

THE FIX IS FREE

Put the prompt and the output in version control.

A prompt someone tuned for four hours is an asset.
So is the finished piece of work it produced.
Both are worth keeping where the team can find them.

Smaller models, your hardware, nothing leaves the building.

AN OPEN-WEIGHT 8B MODEL FROM HUGGING FACE, IN MEMORY

FP16 full precision



INT8 effectively lossless



INT4 quantized further




4 GB runs on a laptop.

Not a data center. Not a contract. A laptop.

Small enough that two or three specialists share one GPU: one for extraction, one for summarization, one for drafting.

CUMULATIVE COST — API vs YOUR OWN HARDWARE
at ~43,000 premium-tier runs a month



And nothing leaves the building.

* Local wins on volume — at commodity cloud pricing, electricity alone can cost more.

What you **don't** put in the prompt is a line on the budget too.

NEVER GOES IN

API keys and passwords

Customer names, emails, card data

Client contracts and their IP

Unreleased financials

FINE TO SEND

The deck you're summarizing

The meeting transcript

Last quarter's public summary


Your own draft email

It costs nothing to strip it out. It costs a great deal to explain it.

Fewer tokens in the prompt is also, literally, a smaller bill. Hygiene and thrift point the same way here.

The most expensive line item is the one that isn't on the invoice.

HOURS FREED PER PERSON, CUMULATIVE



SHARE OF U.S. ADULTS — MEN vs WOMEN

men women

Have used an AI chatbot



Use one every day



Use one for work



Say it makes them more productive



The distance compounds.

“The women who get ahead will be the ones using AI to multiply their output, not just speed up their typing.”

— Ashika Schroll

The appetite is already there. The training is the missing line item.

AI IS THE #1 TOPIC WOMEN SAY THEY WANT TO LEARN

report a lack of skills and access
to training on the job.

ACCESS TO AI TRAINING AT WORK


37% have it

63% do not

The constraint is not appetite. It is access.

WHAT THE C-SUITE ITSELF EXPECTS, WITHIN THREE YEARS

of workers will need reskilling
because of AI.



Nearly half the people you already employ.

Every lever in this talk needs somebody who knows to pull it.

Routing, context and scripting do not happen because you bought a tool. They happen because someone was taught what to look for.

Four moves you can make **without a budget request.**

01 · METER IT

You cannot manage what you're not measuring.

Ask for cost per team, per month — and keep building your intuition for cost per workflow.

02 · ROUTE IT

Name your three biggest use cases. Work through two of them using the cheaper model.

Run the tasks on the cheaper model for two weeks and compare the output. If no material difference can be discerned, you've found savings.

03 · AUTOMATE ONE THING

Pick the task that runs the same way every week.

The Friday report nobody enjoys building. Ask for it to be automated once, and for the recipe to live somewhere the team can find it instead of one person's inbox.

04 · BUDGET PER OUTCOME

Set a price per **thing, not per month.**

A monthly cap tells you when you're in trouble; tracking cost per task tells you whether it was worth it.

This is what “it reads constantly” looks like.



Video plays in the live deck only —
see upskilled.consulting/nextup-socal

25x

READS FOR EVERY WRITE

Reads (in) 500,000/day



Writes (out) 20,000/day



Same shape as Scenario 2, a few slides back. This is what that ratio looks like actually running.

— WRAPPING UP

The takeaways worth keeping.

Tokens are cheap. Volume drives the bill.

Summarizing a hundred-page report costs about a stamp — that alone won't break a budget. Agents that ingest your business artifacts all day and reason in circles about what to do with them can cost far more than anticipated.

Agents read far more than they write.

As we demonstrated, always-on agents can easily cost 47x more than you're accustomed to paying for ad hoc chatbot sessions.

Roughly three-quarters of cost is decision-driven, not a set price.

Routing, context and scripting get you most of the way back — without removing capabilities.

Ask for cost per outcome, not cost per month.

Continue building intuition around cost per task and partnering across departments to begin analyzing monthly spend by task and workflow.

Questions?

Nick McCarty

Upskilled Consulting

nick@upskilled.consulting

linkedin.com/in/nicholasmccarty



upskilled.consulting/nextup-socal/

Scan for the cost checklist, the slides,
and how to reach me.